

COMP3430 / COMP8430

Data wrangling

Lecture 16: Record pair comparison (2)
(Lecturer: Peter Christen)



Lecture outline

- Comparing strings for record linkage
- Approximate string comparison functions
 - Q-gram based (Jaccard and Dice coefficient)
 - Edit and bag distances
 - Jaro-Winkler
- Statistical linkage key SLK-581

Comparing strings for record linkage

- Strings (text) are the most commonly used attributes (fields) when comparing records
 - Names: Title; first, middle, and last name; name suffix and prefix, etc.
 - Addresses: Street name and type, postcode/zipcode (e.g. in the UK: “CB3 0EH”), suburb/town name, state/territory and country names
 - Telephone numbers, emails, credit card numbers, drivers license numbers, etc.
- The aim is to calculate a normalised similarity between two strings with $0 \leq sim_{approx} \leq 1$
- Many different techniques available
 - Some general to any types of strings, others specific to certain types (such as personal names or long genome sequences)

Q-gram based string comparison (1)

- Convert a string into its set of q-grams
 - Often with $q = 2$ (bigrams) or $q = 3$ (trigrams)
 - For example, with bigrams: “peter” \rightarrow [“pe”, “et”, “te”, “er”]
 - Calculate the similarity between two strings based on counting the number of q-grams that occur in both strings
 - Jaccard similarity: $sim_{Jacc}(s_1, s_2) = |intersection(Q_1, Q_2)| / |union(Q_1, Q_2)|$
 - Dice coefficient: $sim_{Dice}(s_1, s_2) = 2 * |intersection(Q_1, Q_2)| / (|Q_1| + |Q_2|)$
- where:
- Q_x is the set of q-grams extracted from string s_x
 - $intersection(Q_1, Q_2)$ is the set of q-grams that occur in both strings
 - $|..|$ denotes the number of elements in a set

Q-gram based string comparison (2)

- For example, with $s_1 = \text{"peter"}$ and $s_2 = \text{"pete"}$ and $q = 2$:
 - $Q_1 = [\text{"pe"}, \text{"et"}, \text{"te"}, \text{"er"}]$, $Q_2 = [\text{"pe"}, \text{"et"}, \text{"te"}]$, $|Q_1| = 4$, $|Q_2| = 3$
 - $\text{intersection}(Q_1, Q_2) = [\text{"pe"}, \text{"et"}, \text{"te"}]$ and $\text{union}(Q_1, Q_2) = [\text{"pe"}, \text{"et"}, \text{"te"}, \text{"er"}]$
 - $\text{sim}_{\text{Jacc}}(s_1, s_2) = |[\text{"pe"}, \text{"et"}, \text{"te"}]| / |[\text{"pe"}, \text{"et"}, \text{"te"}, \text{"er"}]| = 3 / 4 = 0.75$
 - $\text{sim}_{\text{Dice}}(s_1, s_2) = 2 * 3 / (3 + 4) = 6 / 7 = 0.857$
- **Questions:** *Which one is correct? Which one is better?*
What are the Jaccard and Dice similarities between $s_1 = \text{"peter"}$ and $s_2 = \text{"pedro"}$ for $q = 1, 2$, and 3 ?

Edit distance (1)

- Idea: Count how many basic *edit operations* are needed to convert one string into another (known as *Levenshtein* edit distance)
 - Insertion of a character: “pete” → “peter”
 - Deletion of a character: “miller” → “miler”
 - Substitution of a character: “smith” → “smyth”
 - Transpositions of two adjacent characters: “sydney” → “sydeny”
(known as *Damerau-Levenshtein* edit distance)
- **Questions:** *What is the Levenshtein edit distance between “peter” and “petra”, and between “gayle” and “gail”?*

Edit distance (2)

- Convert an edit distance into a similarity $0 \leq sim_{edit_dist} \leq 1$ by calculating $sim_{edit_dist}(s_1, s_2) = 1 - edit_dist(s_1, s_2) / \max(len(s_1), len(s_2))$
- For example, with $s_1 = \text{"peter"}$ and $s_2 = \text{"petra"}$:
$$sim_{edit_dist}(s_1, s_2) = 1 - 2 / \max(5, 5) = 1 - 2 / 5 = 3 / 5 = 0.6$$
- Edit distance can be calculated using a dynamic programming algorithm based on the *edit matrix*
 - Which has a quadratic complexity in the lengths of the two strings (i.e. requires $len(s_1) * len(s_2)$ computational steps)

Edit distance (3)

- Matrix shows the number of edits between sub-strings
(for example, between “ga” and “gayle” → 3 inserts)

“gail” → substitute “i” with “y”,
then insert “e” → “gayle”
(final edit distance is 2)

- Question:** Calculate edit distance
between $s_1 = \text{“peter”}$ and
 $s_2 = \text{“petra”}$

		g	a	y	l	e
	0	1	2	3	4	5
g	1	0	1	2	3	4
a	2	1	0	1	2	3
i	3	2	1	1	2	3
l	4	3	2	2	1	2

Bag distance

- Main drawback of edit distance is its quadratic complexity in the lengths of the two strings, i.e. $len(s_1) * len(s_2)$ computational steps
- A fast approximation of edit distance is *bag distance*
 - A bag is a multi set of the characters in a string: “peter” \rightarrow [“e”, “e”, “p”, “r”, “t”]
- Bag distance is defined as:
 $bag_dist(s_1, s_2) = \max(|x - y|, |y - x|)$, where $x = bag(s_1)$ and $y = bag(s_2)$
- It has been shown that always: $bag_dist(s_1, s_2) \leq edit_dist(s_1, s_2)$, and therefore: $sim_{bag_dist}(s_1, s_2) \geq sim_{edit_dist}(s_1, s_2)$
 - If $sim_{bag_dist}(s_1, s_2)$ is below a threshold then edit distance does not need to be calculated

Jaro-Winkler string comparison (1)

- Developed by the US Census Bureau specifically to compare personal name strings, taking various heuristics into account that are based on extensive practical experiences of name matching
- A combination of q-gram and edit distance string comparison
- Basic idea of the Jaro comparison function:
 - Count c , the number of agreeing (common) characters within half the length of the longer string
 - Count t , the number of transposed characters ('pe' versus 'ep') in the set of common strings
 - Calculate $sim_{Jaro}(s_1, s_2) = (c / len(s_1) + c / len(s_2) + (c - t) / c) / 3$

Jaro-Winkler string comparison (2)

- Further modifications, named Jaro-Winkler, aim to improve name matching further
 - Increase similarity if the first few characters (p , with $p \leq 4$) are the same:
$$\text{sim}_{\text{Jaro_Winkler}}(s_1, s_2) = \text{sim}_{\text{Jaro}}(s_1, s_2) + (1 - \text{sim}_{\text{Jaro}}(s_1, s_2)) * p/10$$
For example, for $s_1 = \text{"peter"}$ and $s_2 = \text{"petra"}$: $p = 3$ ("pet")
 - Further increase the similarity if both strings are at least 5 characters long and contain two common characters besides the prefix
 - Adjust similarity if certain similar character pairs (like in Soundex) occur in two strings (for example "w" \leftrightarrow "v" or "s" \leftrightarrow "z")

Statistical linkage key SLK-581 (1)

- Developed by the Australian Institute for Health and Welfare
<http://meteor.aihw.gov.au/content/index.phtml/itemId/349895>
- Aims to identify records that likely correspond to the same person
- Combines blocking and comparison functionalities
- Basic idea:
 - Take the 2nd, 3rd, and 5th letters of a record's family name (surname)
 - Take the 2nd and 3rd letters of the record's given name (first name)
 - Take the day, month and year of the person, concatenated in that order (*ddmmyyyy*) to form the date of birth
 - Take the gender of the person (1=male, 2=female, 9=unknown)
 - If names too short use 2, if full name component missing use 999

Statistical linkage key SLK-581 (2)

- Examples: (spaces added for illustration only)
 - “marie miller”, 13/04/1991, “f” → “ile ar 13041991 2”
 - “john smith”, 31/03/2001, “m” → “mih oh 31032001 1”
 - “ashley lee”, 11/12/1963, “u” → “ee2 sh 11121963 9”
- **Question:** *Calculate SLK-581 for yourself*